

Avocado Dataset (Project 3)

Name = Chandrima Chatterjee

Problem Statement for the Dataset

"The Avocado dataset we are classifying Organic & Conventional Type and predicting the Average price using Regression model from year 2015, 2016, 2017 and 2018 data."

This Dataset includes the data of consumption of the Avocado fruit in different city of the USA ranging from years from 2015 to 2018.

We have two types of Avocado available here:

1. Organic which is healthy
2. Conventional

The variables on this dataset available are as follows:


1. Categorical: 'region', 'type'
2. Date: 'Date'
3. Numerical: 'Total Volume', '4046', '4225', '4770', 'Total Bags', 'Small Bags', 'Large Bags', 'XLarge Bags', 'Year'
4. Target: 'AveragePrice'

The below dataset is extracted from the different outlets which includes Grocery, mass, clubs, drug, dollar, military units as we can see that the Avocado's are being sold in small to large bags.

The Average Price (of avocados) in the table reflects a per unit cost (per avocado), even when multiple units (avocados) are sold in bags. The Product Lookup codes (PLU's) in the table are only for Hass avocados. Other varieties of avocados (e.g. greenskins) are not included in this table.

Data Analysis

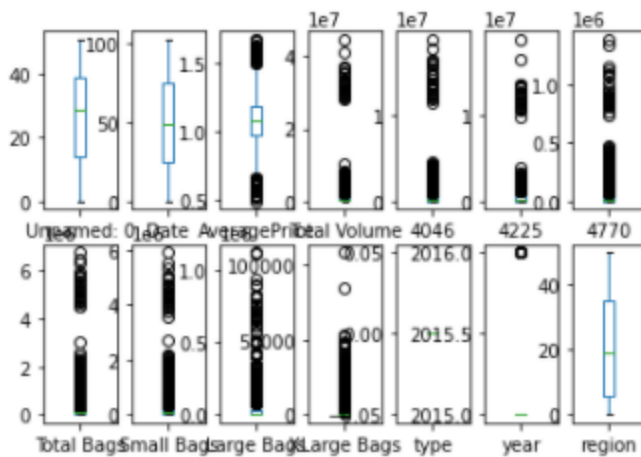
```
] : sns.countplot(x='region',data=df)
plt.show()
```



Target/dependent variable is discrete and categorical in nature -- highest count of region is of california i.e 76. -- the no. of counts of region ranges from 0 to 80. -- lowest count of region is of LosAngles i.e 3. #Lets use LabelEncoder to convert all categorical data into numerical data, so that EDA could be done properly to understand the dataset better.


```
: df.plot(kind='box',subplots = True,layout=(2,7))
```

```
: Unnamed: 0      AxesSubplot(0.125,0.536818;0.0945122x0.343182)
Date             AxesSubplot(0.238415,0.536818;0.0945122x0.343182)
AveragePrice     AxesSubplot(0.351829,0.536818;0.0945122x0.343182)
Total Volume     AxesSubplot(0.465244,0.536818;0.0945122x0.343182)
4046             AxesSubplot(0.578659,0.536818;0.0945122x0.343182)
4225             AxesSubplot(0.692073,0.536818;0.0945122x0.343182)
4770             AxesSubplot(0.805488,0.536818;0.0945122x0.343182)
Total Bags      AxesSubplot(0.125,0.125;0.0945122x0.343182)
Small Bags      AxesSubplot(0.238415,0.125;0.0945122x0.343182)
Large Bags      AxesSubplot(0.351829,0.125;0.0945122x0.343182)
XLarge Bags     AxesSubplot(0.465244,0.125;0.0945122x0.343182)
type            AxesSubplot(0.578659,0.125;0.0945122x0.343182)
year            AxesSubplot(0.692073,0.125;0.0945122x0.343182)
region          AxesSubplot(0.805488,0.125;0.0945122x0.343182)
dtype: object
```




```
: df['Total Volume'].plot.hist()
```

```
: <matplotlib.axes._subplots.AxesSubplot at 0x1f21323bac0>
```




```
: #Checking correlation with the help of heatmap.
sns.heatmap(dfcor)
```

```
: <matplotlib.axes._subplots.AxesSubplot at 0x1f21b584400>
```




```
: #Lets visualize outliers through boxplots
plt.figure(figsize=(ncol,5*ncol))
for i in range (1,len(collist)):
    plt.subplot(nrows,ncol,i+1)
    sns.boxplot(df[collist[i]],color='green',orient='v')
plt.tight_layout()
```



Activate Windows
Go to Settings to

```
plt.figure(figsize=(20,20))
for i in range(0,len(collist)):
    plt.subplot(nrows,ncol,i+1)
    sns.distplot(df[collist[i]])
```



```
: #lets again check the skewness
df.skew()
```

```
: Unnamed: 0      -0.234824
   Date           0.012623
   AveragePrice   -0.109444
   Total Volume    0.442493
   4046           -0.160268
   4225            0.184436
   4770           -0.355508
   Total Bags     0.695502
   Small Bags     0.713843
   Large Bags     -0.912766
   XLarge Bags    0.783913
   year           1.828332
   region         0.288146
   dtype: float64
```

Preprocessing pipeline

```
: #Now treating the outliers
from scipy.stats import zscore
z = np.abs(zscore(df))
z
```

```
: array([[1.81868039, 1.37776563, 1.35048079, ..., 0.81077519, 0.44100815,
          1.3143384 ],
        [1.75131034, 0.57857991, 1.45639674, ..., 0.81077519, 0.44100815,
          1.3143384 ],
        [1.6839403 , 0.22060582, 0.76783831, ..., 0.81077519, 0.44100815,
          1.3143384 ],
        ...,
        [1.01023983, 1.51928262, 2.14485045, ..., 1.10389091, 2.26753179,
          0.88028586],
        [0.94286978, 1.07807099, 2.09189247, ..., 0.81077519, 2.26753179,
          0.88028586],
        [0.87549974, 0.27888526, 1.88006056, ..., 0.81077519, 2.26753179,
          0.88028586]])
```

```
: threshold = 3
print(np.where(z>3))

(array([ 760, 1182, 1182, 1183, 1183, 1184, 1184, 1185, 1185, 1186, 1186,
        1187, 1188, 1188, 1189, 1191, 1346, 1411, 1457, 1458], dtype=int64), array([2, 7, 8, 7, 8, 7, 8, 7, 8, 8, 7, 8, 7, 6, 2,
        2, 2],
        dtype=int64))
```

```
: df_new=df[((z<3).all(axis=1))] #Removing the outliers
```

```
: z[760][2]
```

```
: 3.097989311954043
```

```
: z[1182][3]
```

```
: 2.761987656202092
```

Building Machine Learning Models

```
: from sklearn.svm import SVC

svc=SVC(kernel="linear", C=1)
svc.fit(train_x,train_y)

predsvc =svc.predict(test_x)

print('actual and predicted value score',accuracy_score(test_y,predsvc))
print(confusion_matrix(test_y,predsvc))
print(classification_report(test_y,predsvc ))
```

actual and predicted value score 0.9461077844311377

```
[[20  0  0 ...  0  0  0]
 [ 0  6  0 ...  0  0  0]
 [ 0  0 15 ...  0  0  0]
 ...
 [ 0  0  0 ...  2  0  0]
 [ 0  0  0 ...  0  5  0]
 [ 0  0  0 ...  0  0  5]]
```

	precision	recall	f1-score	support
0	0.95	1.00	0.98	20
1	0.60	0.86	0.71	7
2	0.94	1.00	0.97	15
3	0.94	1.00	0.97	16
4	0.92	1.00	0.96	11
5	1.00	0.89	0.94	9
6	1.00	1.00	1.00	13
7	0.82	0.90	0.86	10
8	1.00	1.00	1.00	5
9	0.89	0.89	0.89	9
10	1.00	0.67	0.80	3
11	1.00	1.00	1.00	4
12	1.00	1.00	1.00	11
13	1.00	1.00	1.00	6
15	1.00	1.00	1.00	1
16	1.00	1.00	1.00	3

```
: from sklearn.neighbors import KNeighborsClassifier
```

```
knn=KNeighborsClassifier(n_neighbors = 5)  
knn.fit(train_x,train_y)
```

```
predknn = knn.predict(test_x)
```

```
print(accuracy_score(predknn,test_y))  
print(confusion_matrix(test_y,predknn))  
print(classification_report(test_y,predknn))
```

```
0.2874251497005988
```

```
[[19  0  0 ...  0  0  0]  
 [ 0  4  0 ...  0  0  0]  
 [ 0  0 11 ...  0  0  0]
```

```
...
```

```
[ 0  0  0 ...  0  0  0]  
[ 0  0  0 ...  0  1  0]  
[ 0  0  0 ...  0  0  0]]
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.51	0.95	0.67	20
1	0.11	0.57	0.18	7
2	0.28	0.73	0.40	15
3	0.39	0.69	0.50	16
4	0.10	0.09	0.10	11
5	0.50	0.33	0.40	9
6	0.62	0.62	0.62	13
7	0.20	0.10	0.13	10
8	0.00	0.00	0.00	5
9	0.30	0.78	0.44	9
10	0.00	0.00	0.00	3
11	0.00	0.00	0.00	4
12	0.29	0.18	0.22	11
13	0.00	0.00	0.00	6
14	0.00	0.00	0.00	0
15	0.00	0.00	0.00	1
16	0.00	0.00	0.00	2

```
: from sklearn.tree import DecisionTreeClassifier
```

```
dct=DecisionTreeClassifier(criterion='entropy')
```

```
dct.fit(train_x,train_y)
```

```
pred_dct=dct.predict(test_x)
```

```
print(accuracy_score(pred_dct,test_y))
```

```
print(confusion_matrix(test_y,pred_dct))
```

```
print(classification_report(test_y,pred_dct))
```

```
0.8652694610778443
```

```
[[18  0  0 ...  0  0  0]
```

```
 [ 0  7  0 ...  0  0  0]
```

```
 [ 0  0 13 ...  0  0  0]
```

```
...
```

```
 [ 0  0  0 ...  2  0  0]
```

```
 [ 0  0  0 ...  0  5  0]
```

```
 [ 0  0  0 ...  0  0  2]]
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.00	0.90	0.95	20
---	------	------	------	----

1	0.58	1.00	0.74	7
---	------	------	------	---

2	0.93	0.87	0.90	15
---	------	------	------	----

3	0.89	1.00	0.94	16
---	------	------	------	----

4	1.00	0.82	0.90	11
---	------	------	------	----

5	0.89	0.89	0.89	9
---	------	------	------	---

6	1.00	0.92	0.96	13
---	------	------	------	----

7	0.67	1.00	0.80	10
---	------	------	------	----

8	1.00	1.00	1.00	5
---	------	------	------	---

9	0.89	0.89	0.89	9
---	------	------	------	---

10	0.00	0.00	0.00	3
----	------	------	------	---

11	0.80	1.00	0.89	4
----	------	------	------	---

12	1.00	0.91	0.95	11
----	------	------	------	----

13	0.86	1.00	0.92	6
----	------	------	------	---

15	0.50	1.00	0.67	1
----	------	------	------	---

16	0.50	0.33	0.40	3
----	------	------	------	---

Conclusion

#From above we can see that svc model has the best score, so we save this model

```
import pickle

# Save to file in the current working directory
pkl_filename = "pickle_model.pkl"
with open(pkl_filename, 'wb') as file:
    pickle.dump(svc, file)

# Load from file
with open(pkl_filename, 'rb') as file:
    pickle_model = pickle.load(file)

# Calculate the accuracy score and predict target values
score = pickle_model.score(test_x, test_y)
print("Test score: {0:.2f} %".format(100 * score))
Ypredict = pickle_model.predict(test_x)
```

Test score: 94.61 %

```
import pickle

filename='picklesvcfile.pkl'
pickle.dump(svc, open(filename, 'wb'))

#load the model from disk

loaded_model=pickle.load(open(filename, 'rb'))

loaded_model.predict(test_x)

array([42, 13, 34, 50, 19, 12, 15, 1, 31, 31, 2, 0, 7, 34, 2, 1, 26,
       4, 41, 9, 5, 9, 31, 6, 11, 12, 17, 22, 20, 1, 44, 12, 37, 3,
       43, 0, 0, 42, 46, 3, 33, 44, 6, 47, 4, 19, 0, 0, 1, 7, 12,
       8, 24, 22, 26, 0, 4, 26, 44, 0, 4, 39, 0, 47, 48, 3, 19, 26,
       31, 34, 38, 47, 2, 4, 0, 34, 38, 46, 26, 18, 47, 50, 26, 44, 34,
       44, 44, 42, 3, 0, 46, 22, 24, 3, 2, 9, 46, 44, 28, 42, 47, 16,
       6, 3, 29, 6, 37, 46, 29, 2, 17, 37, 10, 26, 12, 33, 12, 12, 26,
       9, 31, 36, 36, 7, 26, 49, 2, 35, 32, 6, 31, 33, 8, 17, 17, 30,
       34, 36, 1, 31, 36, 31, 7, 6, 37, 12, 0, 5, 21, 31, 44, 32, 7,
       44, 6, 27, 5, 12, 0, 9, 3, 2, 3, 4, 50, 9, 11, 4, 3, 40,
       0, 33, 33, 42, 49, 25, 6, 19, 42, 49, 17, 26, 7, 25, 44, 43, 37,
       11, 9, 3, 1, 0, 4, 0, 2, 34, 4, 47, 9, 3, 25, 28, 45, 17,
       19, 31, 32, 2, 28, 12, 40, 17, 6, 17, 7, 13, 50, 1, 37, 1, 2,
       44, 6, 23, 29, 31, 2, 6, 3, 41, 6, 0, 0, 0, 1, 7, 20, 45,
       25, 5, 32, 44, 12, 4, 5, 25, 22, 49, 47, 44, 45, 0, 39, 3, 31,
       34, 41, 48, 3, 18, 47, 8, 7, 28, 0, 31, 3, 9, 4, 3, 17, 44,
       47, 2, 18, 49, 18, 4, 25, 7, 2, 50, 36, 31, 29, 5, 38, 17, 32,
       34, 16, 2, 26, 0, 43, 17, 26, 16, 38, 13, 13, 1, 44, 7, 5, 22,
       35, 19, 38, 17, 5, 49, 19, 38, 10, 3, 47, 8, 11, 2, 23, 26, 8,
       39, 13, 47, 33, 13, 42, 2, 47, 23, 6, 38])
```

SVC is the best model for performance where I got **94.61%** of accuracy among of the entire performing machine learning model.